

In [1]:

```
pip install --upgrade pip
```

Requirement already satisfied: pip in c:\users\sowmika\appdata\local\programs\python\python310\lib\site-packages (23.1.2) Note: you may need to restart the kernel to use updated packages.

In [2]:

```
pip install seaborn
```

Requirement already satisfied: seaborn in c:\users\sowmika\appdata\local\programs\python\python310\lib\site-packages (0.12.2)

Requirement already satisfied: numpy!=1.24.0,>=1.17 in c:\users\sowmika\appdata\local\programs\python\python310\lib\site-packages (from seaborn) (1.24.3)

Requirement already satisfied: pandas>=0.25 in c:\users\sowmika\appdata\local\programs\python\python310\lib\site-packages (from seaborn) (2.0.1)

Requirement already satisfied: matplotlib!=3.6.1,>=3.1 in c:\users\sowmika\appdata\local\programs\python\python310\lib\site-packages (from seaborn) (3.7.1)

Requirement already satisfied: contourpy>=1.0.1 in c:\users\sowmika\appdata\local\programs\python\python310\lib\site-packages (from matplotlib!=3.6.1,>=3.1->seaborn) (1.0.7)

Requirement already satisfied: cycler>=0.10 in c:\users\sowmika\appdata\local\programs\python\python310\lib\site-packages (from matplotlib!=3.6.1,>=3.1->seaborn) (0.11.0)

Requirement already satisfied: fonttools>=4.22.0 in c:\users\sowmika\appdata\local\programs\python\python310\lib\site-packages (from matplotlib!=3.6.1,>=3.1->seaborn) (4.39.4)

Requirement already satisfied: kiwisolver>=1.0.1 in c:\users\sowmika\appdata\local\programs\python\python310\lib\site-packages (from matplotlib!=3.6.1,>=3.1->seaborn) (1.4.4)

Requirement already satisfied: packaging>=20.0 in c:\users\sowmika\appdata\local\programs\python\python310\lib\site-packages (from matplotlib!=3.6.1,>=3.1->seaborn) (23.1)

Requirement already satisfied: pillow>=6.2.0 in c:\users\sowmika\appdata\local\programs\python\python310\lib\site-packages (from matplotlib!=3.6.1,>=3.1->seaborn) (9.5.0)

Requirement already satisfied: pyparsing>=2.3.1 in c:\users\sowmika\appdata\local\programs\python\python310\lib\site-packages (from matplotlib!=3.6.1,>=3.1->seaborn) (3.0.9)

Requirement already satisfied: python-dateutil>=2.7 in c:\users\sowmika\appdata\local\programs\python\python310\lib\site-packages (from matplotlib!=3.6.1,>=3.1->seaborn) (2.8.2)

Requirement already satisfied: pytz>=2020.1 in c:\users\sowmika\appdata\local\programs\python\python310\lib\site-packages (from pandas>=0.25->seaborn) (2023.3)

Requirement already satisfied: tzdata>=2022.1 in c:\users\sowmika\appdata\local\programs\python\python310\lib\site-packages (from pandas>=0.25->seaborn) (2023.3)

Requirement already satisfied: six>=1.5 in c:\users\sowmika\appdata\local\programs\python\python310\lib\site-packages (from python-dateutil>=2.7->matplotlib!=3.6.1,>=3.1->seaborn) (1.16.0)

Note: you may need to restart the kernel to use updated packages.

In [3]:

```
!pip install scikit-learn
```

Requirement already satisfied: scikit-learn in c:\users\sowmika\appdata\local\programs\python\python310\lib\site-packages (1.2.2)

Requirement already satisfied: numpy>=1.17.3 in c:\users\sowmika\appdata\local\programs\python\python310\lib\site-packages (from scikit-learn) (1.24.3)

Requirement already satisfied: scipy>=1.3.2 in c:\users\sowmika\appdata\local\programs\python\python310\lib\site-packages (from scikit-learn) (1.10.1)

Requirement already satisfied: joblib>=1.1.1 in c:\users\sowmika\appdata\local\programs\python\python310\lib\site-packages (from scikit-learn) (1.2.0)

Requirement already satisfied: threadpoolctl>=2.0.0 in c:\users\sowmika\appdata\local\programs\python\python310\lib\site-packages (from scikit-learn) (3.1.0)

In [4]:

```
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn import preprocessing, svm
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
```

In [5]:

```
df=pd.read_csv(r"C:\Users\sowmika\Downloads\bottle.csv.zip")
df
```

C:\Users\sowmika\AppData\Local\Temp\ipykernel_19392\259309552.py:1: DtypeWarning: Columns (47,73) have mixed types. Specify dtype option on import or set low_memory=False.

```
df=pd.read_csv(r"C:\Users\sowmika\Downloads\bottle.csv.zip")
```

Out[5]:

Stn_Cnt	Sta_ID	Depth_ID	Depthm	T_degC	Salnty	O2ml_L	STheta	O2Sat	...	R_PHAEO
1	054.0 056.0	19-4903CR-HY-060-0930-05400560-0000A-3	0	10.500	33.4400	NaN	25.64900	NaN	...	NaN
2	054.0 056.0	19-4903CR-HY-060-0930-05400560-0008A-3	8	10.460	33.4400	NaN	25.65600	NaN	...	NaN

In [6]:

```
df=df[['Salnty','T_degC']]
df.columns=['sal','Temp']
```

In [7]:

```
df.head(10)
```

Out[7]:

	sal	Temp								
0	33.440	10.500	19-4903CR-0903CR-HY-060-0930-05400560-0010A-7	19	10.450	33.4200	NaN	25.64300	NaN	NaN
1	33.440	10.460	20-1611SR-0930-05400560-0020A-7	20	10.450	33.4210	NaN	25.64300	NaN	NaN
2	33.437	10.460	20-1611SR-0930-05400560-0020A-7
3	33.420	10.450	20-1611SR-0930-05400560-0020A-7
4	33.421	10.450	20-1611SR-0930-05400560-0020A-7
5	33.431	10.450	20-1611SR-0930-05400560-0020A-7
6	33.440	10.450	0 18.744 33.4083 5.805 23.87055 108.74 ... 0.18	0	18.744	33.4083	5.805	23.87055	108.74	0.18
7	33.424	10.240	20-1611SR-0930-05400560-0020A-7
8	33.420	10.060	20-1611SR-0930-05400560-0020A-7
9	33.431	9.860	2 18.744 33.4083 5.805 23.87072 108.74 ... 0.18	2	18.744	33.4083	5.805	23.87072	108.74	0.18
34859	093.4	026.4	20-1611SR-0930-05400560-0020A-7
34860	093.4	026.4	20-1611SR-0930-05400560-0020A-7
34861	093.4	026.4	5 18.692 33.4150 5.796 23.88911 108.46 ... 0.18	5	18.692	33.4150	5.796	23.88911	108.46	0.18
34862	093.4	026.4	10 18.161 33.4062 5.816 24.01426 107.74 ... 0.31	10	18.161	33.4062	5.816	24.01426	107.74	0.31


In [8]:

```
df_Cnt Sta_ID Depth_ID Depthm T_degC Salnty O2ml_L STheta O2Sat ... R_PHAEO
```

```
sns.lmplot(x="sal", y="Temp", data=df, order=2, ci=None)
```

Out[8]:

```
1611SR-
093.4 MX-310-
34863 026.4 1289-
<seaborn.axisgrid.FacetGrid at 0x1f5dceae950>
```



In [9]:

```
df.describe()
```

Out[9]:

	sal	Temp
count	817509.000000	853900.000000
mean	33.840350	10.799677
std	0.461843	4.243825
min	28.431000	1.440000
25%	33.488000	7.680000
50%	33.863000	10.060000
75%	34.196900	13.880000
max	37.034000	31.140000

In [10]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 864863 entries, 0 to 864862
Data columns (total 2 columns):
#   Column  Non-Null Count  Dtype
---  -
0    sal      817509 non-null    float64
1   Temp      853900 non-null    float64
dtypes: float64(2)
memory usage: 13.2 MB
```

In [29]:

```
df.fillna(method = 'ffill',inplace = True)
```

C:\Users\sowmika\AppData\Local\Temp\ipykernel_19392\3028625988.py:1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
df.fillna(method = 'ffill',inplace = True)
```

In [12]:

```
X = np. array(df['sal']).reshape(-1, 1)
Y = np.array(df['Temp']).reshape(-1, 1)
```

In [30]:

```
df.dropna(inplace = True)
```

C:\Users\sowmika\AppData\Local\Temp\ipykernel_19392\1791587065.py:1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
df.dropna(inplace = True)
```


In [14]:

```
X_train,X_test,Y_train,Y_test=train_test_split(X,Y,test_size=0.25)
regr = LinearRegression()
regr.fit(X_train, Y_train)
print(regr.score(X_test, Y_test))
```

0.2036387277835181


In [15]:

```
y_pred=regr.predict(X_test)
plt.scatter(X_test, Y_test,color='b')
plt.plot(X_test,y_pred,color='k')
plt.show()
```



In [16]:


```
df500 = df[:][:500]
sns.lmplot(x ="sal",y ="Temp", data = df500,order = 1,ci = None)
```



In [17]:

```
df500.fillna(method = 'ffill',inplace = True)
X = np. array(df500['sal']).reshape(-1, 1)
y = np.array(df500['Temp']).reshape(-1, 1)
df500.dropna(inplace = True)
X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.25)
regr = LinearRegression()
regr.fit(X_train,y_train)
print("Regression:",regr.score(X_test,y_test))
y_pred=regr.predict(X_test)
plt.scatter(X_test,y_test,color = 'b')
plt.plot(X_test,y_pred,color = 'k')
plt.show()
```

Regression: 0.8429543309817631



In [18]:

```
from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score
model = LinearRegression()
model.fit(X_train,y_train)
y_pred = model.predict(X_test)
r2=r2_score(y_test,y_pred)
print("R2 score:",r2)
```

R2 score: 0.8429543309817631

vehicle selection

In [19]:

```
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn import preprocessing, svm
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
```

In [47]:

```
df=pd.read_csv(r"C:\Users\sowmika\Downloads\fiat500_VehicleSelection_Dataset (1).csv")
df
```

Out[47]:

	ID	model	engine_power	age_in_days	km	previous_owners	lat	
0	1	lounge	51	882	25000	1	44.907242	8.611
1	2	pop	51	1186	32500	1	45.666359	12.241
2	3	sport	74	4658	142228	1	45.503300	11.417
3	4	lounge	51	2739	160000	1	40.633171	17.634
4	5	pop	73	3074	106880	1	41.903221	12.495
...
1533	1534	sport	51	3712	115280	1	45.069679	7.704
1534	1535	lounge	74	3835	112000	1	45.845692	8.666
1535	1536	pop	51	2223	60457	1	45.481541	9.413
1536	1537	lounge	51	2557	80750	1	45.000702	7.682
1537	1538	pop	51	1766	54276	1	40.323410	17.568

1538 rows × 9 columns

In [48]:

```
df=df[['age_in_days','km']]
df.columns=['age','km']
```

In [49]:

```
df.head(10)
```

Out[49]:


	age	km
0	882	25000
1	1186	32500
2	4658	142228
3	2739	160000
4	3074	106880
5	3623	70225
6	731	11600
7	1521	49076
8	4049	76000
9	3653	89000

In [50]:

```
sns.lmplot(x="age", y="km", data=df, order=2, ci=None)
```

Out[50]:

<seaborn.axisgrid.FacetGrid at 0x1f58a24eaa0>



In [51]:

```
df.describe()
```

Out[51]:

	age	km
count	1538.000000	1538.000000
mean	1650.980494	53396.011704
std	1289.522278	40046.830723
min	366.000000	1232.000000
25%	670.000000	20006.250000
50%	1035.000000	39031.000000
75%	2616.000000	79667.750000
max	4658.000000	235000.000000

In [52]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1538 entries, 0 to 1537
Data columns (total 2 columns):
#   Column  Non-Null Count  Dtype
---  -
0   age      1538 non-null    int64
1   km       1538 non-null    int64
dtypes: int64(2)
memory usage: 24.2 KB
```

In [53]:

```
df.fillna(method = 'ffill', inplace = True)
```

C:\Users\sowmika\AppData\Local\Temp\ipykernel_19392\3028625988.py:1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
df.fillna(method = 'ffill', inplace = True)
```

In [54]:

```
X = np.array(df['age']).reshape(-1, 1)
Y = np.array(df['km']).reshape(-1, 1)
```

In [55]:

```
df.dropna(inplace = True)
```

C:\Users\sowmika\AppData\Local\Temp\ipykernel_19392\1791587065.py:1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
df.dropna(inplace = True)
```


In [56]:

```
X_train,X_test,Y_train,Y_test=train_test_split(X,Y,test_size=0.25)
regr = LinearRegression()y_pred=regr.predict(X_test)
plt.scatter(X_test, Y_test,color='b')
plt.plot(X_test,y_pred,color='k')
plt.show()
regr.fit(X_train, Y_train)
print(regr.score(X_test, Y_test))
```

0.6885339929552072

In [32]:

```
y_pred=regr.predict(X_test)
plt.scatter(X_test, Y_test,color='b')
plt.plot(X_test,y_pred,color='k')
plt.show()
```




In [61]:

```
df500 = df[:][:500]  
sns.lmplot(x="age", y="km", data=df500, order=1, ci=None)
```

Out[61]:


<seaborn.axisgrid.FacetGrid at 0x1f58a516fb0>



In [62]:

```
df500.fillna(method = 'ffill',inplace = True)
X = np. array(df500['age']).reshape(-1, 1)
y = np.array(df500['km']).reshape(-1, 1)
df500.dropna(inplace = True)
X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.25)
regr = LinearRegression()
regr.fit(X_train,y_train)
print("Regression:",regr.score(X_test,y_test))
y_pred=regr.predict(X_test)
plt.scatter(X_test,y_test,color = 'b')
plt.plot(X_test,y_pred,color = 'k')
plt.show()
```

Regression: 0.7069135527596802



In [59]:

```
from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score
model = LinearRegression()
model.fit(X_train,y_train)
y_pred = model.predict(X_test)
r2=r2_score(y_test,y_pred)
print("R2 score:",r2)
```

R2 score: 0.7670566623468238

House price prediction

In [63]:

```
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn import preprocessing, svm
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
```

In [64]:

```
df=pd.read_csv(r"C:\Users\sowmika\Downloads\data.csv")
df
```

Out[64]:

	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront
0	2014-05-02 00:00:00	3.130000e+05	3.0	1.50	1340	7912	1.5	0
1	2014-05-02 00:00:00	2.384000e+06	5.0	2.50	3650	9050	2.0	0
2	2014-05-02 00:00:00	3.420000e+05	3.0	2.00	1930	11947	1.0	0
3	2014-05-02 00:00:00	4.200000e+05	3.0	2.25	2000	8030	1.0	0
4	2014-05-02 00:00:00	5.500000e+05	4.0	2.50	1940	10500	1.0	0
...
4595	2014-07-09 00:00:00	3.081667e+05	3.0	1.75	1510	6360	1.0	0
4596	2014-07-09 00:00:00	5.343333e+05	3.0	2.50	1460	7573	2.0	0
4597	2014-07-09 00:00:00	4.169042e+05	3.0	2.50	3010	7014	2.0	0
4598	2014-07-10 00:00:00	2.034000e+05	4.0	2.00	2090	6630	1.0	0
4599	2014-07-10 00:00:00	2.206000e+05	3.0	2.50	1490	8102	2.0	0

4600 rows × 18 columns



In [65]:

```
df=df[['sqft_living','sqft_lot']]  
df.columns=['living','lot']
```

In [66]:

```
df.head(10)
```

Out[66]:


	living	lot
0	1340	7912
1	3650	9050
2	1930	11947
3	2000	8030
4	1940	10500
5	880	6380
6	1350	2560
7	2710	35868
8	2430	88426
9	1520	6200

In [67]:

```
sns.lmplot(x="living", y="lot", data=df, order=2, ci=None)
```

Out[67]:

<seaborn.axisgrid.FacetGrid at 0x1f58a758f70>



In [69]:

```
df.describe()
```

Out[69]:

	living	lot
count	4600.000000	4.600000e+03
mean	2139.346957	1.485252e+04
std	963.206916	3.588444e+04
min	370.000000	6.380000e+02
25%	1460.000000	5.000750e+03
50%	1980.000000	7.683000e+03
75%	2620.000000	1.100125e+04
max	13540.000000	1.074218e+06

In [70]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4600 entries, 0 to 4599
Data columns (total 2 columns):
#   Column  Non-Null Count  Dtype
---  -
0   living   4600 non-null     int64
1   lot      4600 non-null     int64
dtypes: int64(2)
memory usage: 72.0 KB
```

In [71]:

```
df.fillna(method = 'ffill',inplace = True)
```

C:\Users\sowmika\AppData\Local\Temp\ipykernel_19392\3028625988.py:1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
df.fillna(method = 'ffill',inplace = True)
```

In [72]:

```
X = np. array(df['living']).reshape(-1, 1)
Y = np.array(df['lot']).reshape(-1, 1)
```


In [73]:

```
X_train,X_test,Y_train,Y_test=train_test_split(X,Y,test_size=0.25)
regr = LinearRegression()
regr.fit(X_train, Y_train)
print(regr.score(X_test, Y_test))
```

0.06680986187711502

In [74]:

```
y_pred=regr.predict(X_test)
plt.scatter(X_test, Y_test,color='b')
plt.plot(X_test,y_pred,color='k')
plt.show()
```




In [75]:

```
df500 = df[:][:500]  
sns.lmplot(x = "living", y = "lot", data = df500, order = 1, ci = None)
```

Out[75]:


<seaborn.axisgrid.FacetGrid at 0x1f58a75a590>



In [79]:

```
df500.fillna(method = 'ffill',inplace = True)
X = np. array(df500['living']).reshape(-1, 1)
y = np.array(df500['lot']).reshape(-1, 1)
df500.dropna(inplace = True)
X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.25)
regr = LinearRegression()
regr.fit(X_train,y_train)
print("Regression:",regr.score(X_test,y_test))
y_pred=regr.predict(X_test)
plt.scatter(X_test,y_test,color = 'b')
plt.plot(X_test,y_pred,color = 'k')
plt.show()
```

Regression: 0.45239880007632904



In [80]:

```
from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score
model = LinearRegression()
model.fit(X_train,y_train)
y_pred = model.predict(X_test)
r2=r2_score(y_test,y_pred)
print("R2 score:",r2)
```

R2 score: 0.45239880007632904

In []:

